

WHITE PAPER I

Zoning the Namespace

Land-use law as the legal technology of admission

Author	S. Harrison Knudson, J.D. — Director, IFC-SiS
Version	2.0 — August 2026 (supersedes v1.0)
Companion	White Paper II: The Admission Layer (technical)
Correspondence	shknudson@ifcsis.org

Abstract

Regulation of children's online safety has been built almost entirely on the assumption that harmful material can be intercepted in transit. The vocabulary gives it away — upstream and downstream, flow, source, filtration — and it describes water treatment. But treatment presumes a bounded set of sources, and on the internet the upstream is universal: anyone may put anything in, from anywhere, under no licence. At the other end there is no geography either — a river reaches an ocean where everything is mixed, refuse reaches a tip where it piles. Nothing can be intercepted because the sources are unbounded, and nothing arrives sorted because the destination is undifferentiated. Every intervention built on the treatment assumption therefore acts after publication, which is after exposure.

This paper argues that the discipline which already solved the analogous problem in physical space is land-use law. Zoning is a century-old legal technology for deciding what may sit next to what: a controlled vocabulary of uses, definitions attached to each, and permissions that follow from the category. It operates in every developed jurisdiction, is administered by a profession, and is not understood as censorship.

The paper's central claim is a transposition rather than an identity. Public health is a substantial part of zoning's genealogy — its nineteenth-century origins lie in separating noxious trades from dwellings, and Euclid reasoned from nuisance, a use harmless in itself becoming a nuisance by reason of its location. The technique zoning developed is the management of externalities through separation, and a children's district applies that technique to a population the police power has always treated as warranting protection. Paediatric medicine reaches a comparable conclusion without geography, and locates the writing of standards with clinical bodies rather than legislatures.

The paper also observes that physical establishments open to the public routinely operate within several overlapping regulatory layers — land-use control, licensing where applicable, display rules in regulated contexts, and duties to those who enter — while an internet domain, an address to which the public is invited, carries none of them. That is an exemption nobody granted deliberately.

It argues in addition that zoning is, in its ordinary operation, deliberately vague, and that the vagueness is the mechanism rather than the defect. Discretion is the asset; the undefined term is the lever. Content governance has the same structure, and an admission standard is the move that removes the lever: the decision follows from a published classification rather than from a later negotiation.

The claim that uniform structure with local content can work has a precedent rather than merely an analogy: model building codes are adopted, with local amendment, almost everywhere, and a shared technical vocabulary makes locally differing rules interoperable. That is the claim the metadata standard makes.

The paper sets out the transposition, examines what does not survive it, distinguishes the resulting architecture from censorship along five specified conditions, grounds the household election in the UN Convention on the Rights of the Child, and identifies the questions on which the argument most invites contest.

1. The Wrong Metaphor

Every developed legal system has spent a century deciding what may sit next to what. A liquor store may not open beside a primary school. Heavy industry does not belong in a residential district. Nobody describes these rules as censorship, because they are not judgments about whether the liquor store may exist. They are judgments about where.

Children have always been protected by distance. Cities have neighbourhoods a child is not taken to. Shops have shelves above reach; cinemas have doors; libraries have sections. Protection has been geographic, which is to say it has depended on the harmful thing being somewhere else and on reaching it taking effort an adult could notice and interrupt.

Point-to-point addressing deleted that. There is no upstream of a hyperlink and no downstream of a URL; the protocol cannot know which of two addresses is source and which is destination. For a child this means schoolwork and a self-harm forum sit in the same address space, one reach apart. Exposure does not require seeking. It requires only adjacency, and everything is adjacent.

Of the environments children inhabit, the internet is the one in which proximity does least protective work — and arguably none.

1.1 What follows from the river assumption

Regulators across jurisdictions have responded with instruments built on the assumption that content moves through a medium which can be cleaned. Intermediaries are cast as treatment plants and required to clean what passes. The Online Safety Act, the Digital Services Act, and the age-verification statutes proliferating at state level in the United States differ in scope and mechanism but share this shape.

The river is real. The assumption about it is not. Every model of water treatment presumes a bounded set of sources: a known number of outfalls, licensed, locatable and answerable. Filtration works because the upstream is finite. On the internet the upstream is universal. Anyone may put anything into it, from anywhere, at no cost and under no licence, and gravity carries it to whoever is downstream. There is no limit on who dumps, and no point through which all of it passes.

A treatment plant needs a river with banks and a countable number of outfalls above it. This one has neither. The limit on filtration is therefore architectural rather than a matter of diligence or funding: there is nothing that all of it must pass through, so no filter can guarantee the composition of what arrives.

And the other end is worse. A river reaches an ocean, where everything that entered it anywhere is mixed together and nothing can be told apart. Refuse reaches a tip, where it accumulates in a single heap with no order at all. Either image will do, and they say the same thing: there is no geography at the end of the flow.

That is where the child is standing. Not beside a river that could be filtered upstream, but at the point where everything that was put in anywhere has arrived together, unsorted, with the schoolwork and the self-harm forum in the same heap. Protection has always depended on geography, and the end of the flow does not have any.

So there are two difficulties rather than one, and they compound. Interception has no chokepoint, because the sources are unbounded. Arrival is unsorted, because the destination is undifferentiated. Filtering can and does work in particular cases; what it cannot do is guarantee the composition of what a child encounters, and improvements in detection generally arrive after publication, which is after exposure.

The interventions are doing what they were designed to do. They subtract from an open namespace, and the namespace may be the wrong unit of intervention.

One inference must be resisted, and this paper tries to resist it throughout. That filtration has identifiable limits is not evidence that admission gating works. The two are different architectures with different failure models, and the failure of one establishes nothing about the success of the other. What follows from the diagnosis is narrower: admission gating fails differently, its failure modes are testable, and it therefore deserves to be tested rather than adopted. Everything in this paper should be read as an argument for the experiment, not as a claim about its result.

THE RESERVOIR

If the river cannot be treated and the ocean cannot be sorted, the alternative is a reservoir: a separate body of drinking water that fills by permission rather than by gravity. Nothing enters because it happened to flow downhill; it enters because it was vetted and admitted, under an agreement that carries terms. Aqueducts carry it to households that elect it. Nothing is drawn from the river, which is left exactly as it is.

2. Filtration Avoidance

The reservoir is a figure, and figures prove nothing on their own. What follows is offered because the same regulatory move — protect the source rather than treat the supply — appears independently across legal traditions and, in one modern case, was accepted by a regulator in place of a statutory treatment requirement.

2.1 The oldest version of the argument

Reservoirs and aqueducts are ancient, and so is the law protecting them. What is striking is how independently the same instrument appears — not one tradition inventing it and the others borrowing, but several arriving separately at the conclusion that a protected supply requires governing what may happen around and above it.

Rome. A decree of the senate in 11 BC established a reserved zone on either side of the aqueducts, with penalties for violation, and a fine was prescribed for polluting the water. The same reforms created the office of curator aquarum, and the lex Quinctia of 9 BC gave that office powers of access and enforcement. That is a protection district around a water supply, defined in law, with sanctions attached to what may be done within it.

India. The Arthashastra prescribes fines for fouling near a water reservoir, and penalties for damaging, obstructing or polluting waterworks serving settlements and cultivated land. Encroachment on public waterworks is punished severely. The provisions are about placement rather than cleaning — what may be done where, relative to the supply.

Israel. Deuteronomy sets aside a place outside the camp and requires that waste be buried there. It is a siting rule rather than a hygiene one: the point is where the thing goes, not how it is dealt with afterwards. Do not put the latrine upstream of the settlement is among the oldest provisions in law, and it is a zoning provision.

Others could be added, and the dating of each is a question for people better placed to answer it than the author of this paper. The claim here is not that any one is first. It is that the pairing of protected supply and governed surroundings is not a modern invention and never was, and that traditions with no contact with one another reached it independently — which is generally a sign that a thing follows from the problem rather than from a culture.

Rome did not propose to treat the water on arrival. Neither did the others. They governed what could happen near the channel, because governing what enters is the only approach that scales when the supply is long and the settlement is large. Reservoir and zoning have always come together, because neither works without the other.

2.2 The modern version, and the regulator who approved it

The Safe Drinking Water Act of 1986 required filtration of every public surface water supply in the United States. New York City draws roughly nine tenths of its water from the Catskill and Delaware watersheds, and the plant required would have cost between eight and ten billion dollars to build, with operating costs of several hundred

million a year. For a city already carrying that capital programme, the figure was described by its own commissioner as catastrophic.

In 1989 the Environmental Protection Agency ruled that the city could avoid filtration if it could meet the standard another way. The city's answer was to control what may enter the watershed. Land acquisition and conservation easements across five counties. Land-use regulation of what may be built, farmed or discharged upstream of the supply. Wastewater plants upgraded, septic systems replaced, stream corridors stabilised, agricultural practice changed.

THE INSTRUMENT HAS A NAME

The EPA's formal determination is called a Filtration Avoidance Determination. It was first granted in 1997, renewed in 2007 and 2017, and revised in 2022.

A city under a legal duty to filter was permitted to discharge that duty by controlling admission to the watershed instead. Not by treating what arrived, but by governing what may enter.

Nearly thirty years on it is working. Around 2.5 billion dollars has been spent on watershed protection, against eight to ten billion in construction and several hundred million a year in operation for the alternative. The system supplies roughly a billion gallons a day to some nine and a half million people, arriving through aqueducts and tunnels running more than 175 miles from upstate forests. It is the largest unfiltered public water supply in the United States.

2.3 Why this is the precedent and not the parallel

The precedent is narrower than identity, and more useful for being so. New York did not control admission to its watershed in the sense this paper proposes for a namespace: people, farms, towns and septic systems remained, and the programme is a package of protection measures rather than a gate. What the case establishes is the regulatory proposition, and it is the one that matters: a regulator accepted sustained upstream source protection as the basis for avoiding an otherwise applicable downstream treatment requirement. Several structural features recur.

- **A reservoir.** A separate body of water, not drawn from the general flow, whose quality is a property of what was permitted into its watershed.
- **Zoning as protection.** Land-use control is not decorative here. It is the mechanism. What may exist upstream of a reservoir is the thing that determines whether the water arriving is safe, and land-use regulation is how that question is answered.
- **Aqueducts.** Delivery to households outside the flow of any river, carrying only what the protected watershed produced.
- **Filtration avoidance.** A regulator accepting sustained upstream source protection in place of downstream treatment.

That last point is the one this paper most wants to press. The objection to an admission-based approach is usually that regulators will not accept anything other than treatment — that duties are discharged by inspecting what passes, and a proposal to govern the source instead is naive about how regulation works. New York is a counter-example, and not a small one. Faced with a filtration mandate, a city proposed sustained upstream source protection, and the federal regulator accepted it, and has continued to accept it for three decades on the evidence.

The economics motivate a question rather than supplying an answer. Source protection there cost a fraction of what treatment would have cost. Whether the same holds for a namespace — whether upstream admission can be maintained at lower marginal cost than continuous downstream classification — is precisely what is not known, and is among the reasons for running a trial rather than asserting a saving.

One caution, stated because it is the honest one. New York achieved this partly through purchase and easement — buying the watershed, or buying the right to restrict it — backed by a city with the capital to do so and a history of using eminent domain that generated real resentment upstate. A namespace cannot buy its watershed and should not want to. What transfers is the structure and the regulatory precedent, not the method of acquisition.

3. Zoning Is Health Law

The previous section may read as a series of apt parallels. It is not. Every example given — the latrine outside the camp, the fine for fouling near the reservoir, the reserved zone beside the aqueduct, the watershed bought rather than filtered — belongs to a single body of law, and modern zoning belongs to it too.

The police power is conventionally stated as the authority to regulate for the health, safety, morals and general welfare of the public. Zoning is that power exercised through placement. Its nineteenth-century origins are unambiguous: noxious trades separated from dwellings, tenement regulation for light and air, slaughterhouses and tanneries moved away from housing. The instrument was invented to manage externalities that were making people ill.

Euclid reasoned from nuisance. A use harmless in itself may become a nuisance by reason of its circumstances or location — the pig in the parlour rather than the barnyard. Not that the pig is wrong; that the placement is. The case supports a proposition sufficient for the argument here: context and placement can change the regulatory significance of an otherwise lawful use, and the remedy may be separation rather than suppression.

The claim here is a transposition, not an identity. Zoning is not simply health law; the police power runs to safety, welfare, planning, amenity and more besides. But public health is a substantial part of its genealogy, and separation of incompatible uses to manage externalities is the technique it developed. A children's district is an application of that technique to a population the police power has always treated as warranting protection. That is a conventional move, not a novel one — which is different from saying the internet is land.

3.1 The same logic without the geography

Medicine reaches the identical conclusion by a different route, and the parallel is worth stating because it explains something zoning cannot.

Paediatric dosing is not a judgment about a drug. Weight-based dosing, age-banded formulations and paediatric contraindication all rest on a single premise: a developing organism processes the same substance differently. Aspirin is not withheld from children because aspirin is bad, but because of Reye's syndrome. Certain antibiotics are avoided in children because of effects on developing cartilage. Nobody describes any of this as discrimination against the compound.

That is the answer to the objection that an admission standard is content regulation wearing a costume. It is content regulation in the sense that dosing is substance regulation — which is to say, not a judgment about the material's worth, but about what a developing system can process. The material is not condemned. The recipient is different.

It also points toward the kind of expertise that should dominate criterion-setting. Paediatric dosing standards are set by clinical and professional bodies; regulators adopt them rather than authoring them, and a legislature setting its own dosing schedule by statute would be understood to have overstepped. The parallel suggests a similar division here — substantive criteria drawn from developmental psychology, paediatrics, child development, education, online-safety research and children's rights, with the state adopting rather than writing. It suggests that division; it does not settle it, and the composition of the body that writes the criteria is a governance question in its own right.

The obvious objection, and it does not have a tidy answer. Pharmacokinetics is measured; the developmental evidence on content is contested and effect sizes are argued fiercely. It is tempting to reply that a protective standard requires only foreseeability rather than settled effect size, and there is something in that — medicine

contraindicates on signal and mechanism long before the epidemiology is complete. But foreseeability is a liability concept and does not by itself establish the evidentiary threshold for excluding material from a namespace.

The honest position is that the threshold is itself a governance question. A protective code cannot require scientific certainty before acting; nor can it convert contested developmental hypotheses into categorical exclusions merely by invoking precaution. The code must therefore specify what quality and quantity of evidence justifies an exclusion, and that specification must be contestable through the published process like any other provision. Section 15 treats this as an open question, because it is one.

4. Discretion, and Why an Admission Standard Trades It Away

A comprehensive plan is supposed to declare in advance what may be built where. In practice much of it is written in aspirational language and discretionary standards, so that the decision on any particular application is taken later, in a hearing, against criteria that leave considerable latitude.

Whether that latitude is a drafting habit, a deliberate preservation of leverage, or simply the price of accommodating circumstances that cannot be anticipated is a contested question in the planning literature and this paper takes no position on it. What matters here is narrower and not seriously disputed: discretionary standards expressed in natural language are difficult to make predictable in advance, and impossible to make machine-readable.

Content governance has the same structure. Terms of service are drafted broadly. Enforcement is discretionary and applied after the fact. Publishers cannot know in advance what is permitted; users cannot predict what will be removed; and no external party can verify a decision against a stated rule, because the stated rule does not determine the decision.

An admission standard trades discretion for predictability, deliberately. A closed vocabulary with defined terms and published criteria gives up the flexibility to handle the unanticipated case gracefully, and buys in exchange a decision that follows from a classification, can be recorded, can be checked by someone who was not present, and can be appealed against a rule rather than against a judgment. That is a real trade with real costs, and Section 15 treats the question of what it costs as open.

5. The Existence Proof

The strongest evidence that uniform structure with local content works is not an argument. It is that the adjacent discipline already did it.

Structural, fire, plumbing, mechanical and electrical requirements are not reinvented nine thousand times. They are published as model codes — the I-Codes maintained by the International Code Council — and adopted, with local amendment, almost everywhere. An engineer, manufacturer or builder can rely on a shared technical language and know that a beam rated in one county is rated in the next.

Highway engineering did the same. Fifty thousand jurisdictions were not permitted to invent private meanings for traffic terms; the vocabulary was standardised and the profession is stronger for it.

Zoning stands nearly alone in insisting its basic nouns be redefined in every town. Setback, frontage, use, lot coverage — concepts identical from coast to coast, expressed as though a foot measured something different in one county than the next. It is worth asking plainly why the one field whose terms are endlessly re-created is also the field where discretion and delay are most valuable to those administering it.

UNIFORM STRUCTURE, LOCAL CONTENT

The safety codes did not tell a jurisdiction what to require. They fixed the vocabulary in which requirements are expressed, and left the substantive thresholds to local determination. A community keeps every ounce of its authority to decide its own character; it simply loses the ability to hide the answer behind bespoke terminology.

That is exactly the claim the metadata standard makes, and it is not a new claim. It is a claim that has already succeeded once, in the discipline next door.

The proposal here is correspondingly narrow. Not a national ordinance, and not a global content policy — nobody will ever write either, and attempting it wastes the effort. A standardised set of defined terms, so that words mean the same thing everywhere even as the local choices differ.

5.1 France: conformance as the condition of legal effect

The second proof is closer still, and it is recent enough that most of the English-language literature has not caught up with it.

France established the Géoportail de l'urbanisme by ordinance in 2013 and, alongside it, the CNIG geostandards — national specifications prescribing the content and structure of digitised local planning documents: zoning categories, prescriptions, annexes. They are versioned, currently at v2024-01, revised through a published working group with a public issues process, and distributed as open data schemas. The substantive planning choices remain entirely local.

THE MECHANISM WORTH BORROWING

Publication of a plan in the standard on the national portal is reported to be a condition of its opposability — its enforceability against an applicant — from 1 January 2020. This paper states the proposition as reported and flags it for primary citation; the precise legal consequence attaching to conformity is a question of French administrative law and should be footnoted from the ordinance and implementing decree rather than summarised.

If the proposition holds, the mechanism is instructive: a non-conforming plan is not unlawful, it simply does not acquire legal effect. Conformance becomes the condition of enforceability rather than a compulsion, and the choice not to conform remains available at the cost of the benefit conformance confers.

That is the enforcement structure this paper proposes, and it is the answer to the objection that an admission standard must ultimately be backed by coercion: a standard can be made to matter by attaching a benefit to conformance rather than a penalty to its absence. What is beyond dispute is the weaker and still useful proposition — that a national standard prescribing structure while leaving substance local has been in operation for years and is maintained through a public process.

The two proofs answer different objections. The model building codes show that uniform structure with local content is adopted voluntarily and at scale. France shows how a standard acquires force without becoming a mandate. Neither required inventing anything.

5.2 Determination, not negotiation

Once the principles of a code are set down in standardised terms, a plan can compute what a parcel is entitled to rather than routing every application through a discretionary review.

Admission becomes a determination rather than a negotiation. The classification happens before the thing needs a decision, and the decision follows from the classification. That sentence describes automated zoning validation and it describes write-gating; they are the same operation performed on different objects.

This is not theoretical, and it is not new. By 2006 jurisdictions in North Carolina — Raleigh and Charlotte-Mecklenburg — had built interlocal agreements for shared plan databases, reciprocal acceptance of residential reviews, and automated validation of building placement against zoning restrictions, with the stated goal of letting builders self-issue permits over the web with site and zoning validation as the last piece to automate. Portland and the State of Oregon were moving the same direction. The engine existed twenty years ago.

What it lacked, and still lacks, is the standardised dictionary underneath it, so that the automation can run anywhere rather than being rebuilt one jurisdiction at a time. The same gap, in the same shape, is what an admission standard for a namespace is proposing to fill.

6. The Transposition

Land-use regulation is not primarily about buildings. It is a technology for attaching permissions to categories, and its components are precisely the components an admission standard requires.

- **A controlled vocabulary.** A zoning code enumerates uses and the list is closed. A use not in the vocabulary must be classified into it before it can be permitted or refused.
- **Definitions attached to each term.** The code does not merely name categories; it defines them, and the definitions do the work when a case is contested.
- **Permissions that follow from the category.** Once a use is classified, what is permitted follows from the classification and the district, not from a fresh judgment about the applicant.
- **Districts with differing thresholds.** The same vocabulary governs everywhere; what varies is which uses each district permits. Structure invariant, thresholds local.
- **Determination at entry.** The decision is made before the use begins, not after complaint.
- **Published criteria and documented findings.** Decisions are made against a published code and recorded with reasons.
- **Appeal.** An applicant refused may contest the refusal before a body other than the one that refused.
- **Enforcement through tenancy, not interception.** A non-conforming use is cited, given a cure period, penalised, and ultimately removed. Nobody stands at the door inspecting.

Each of these transposes. A namespace admitting registrations against published criteria, recording the decision in machine-readable form, calibrating thresholds by jurisdiction, and enforcing through terms of registration rather than traffic inspection, has the structure of a land-use regime. The correspondences, stated carefully, are these.

Land-use system	This architecture
Model code body	The Foundation, which maintains the invariant schema
Planning authority setting local rules	The national policy body, which sets the applicable admission floor
Zoning or building reviewer	The accredited conformity reviewer, who applies the published criteria
Land records and permit administration	The registry operator, which records and attests operative status
Applicant or property holder	The registrant
Permit status and continuing conformity	The conforming delegation and its tenancy

Appeal body	Independent review, separate from the deciding body
District	The governed namespace

The registry is not the planning authority. It administers the record and executes remedies; it does not set the code or make the substantive judgment. Conflating those roles would put an operator in the position of certifying content suitability, which is precisely the burden this architecture exists to spare them, and would contradict the governance condition in Section 9.

6.1 Districts have internal grammars

Single-use zoning is not only about keeping incompatible uses apart. It recognises that different uses have different internal logics, and that the regulatory framework must respect those logics for the uses to function.

A school is not merely a building from which harmful things are excluded. It has classrooms, a playground, a library, and a particular kind of adult whose role is defined by what they are there to do with the children present. An amusement park is not a non-school; it has its own grammar of queues, rides, vendors and visible supervision. A cinema is different again — seating, darkness, an expected stillness. Children move between these environments, and the transitions themselves teach what kind of behaviour, trust and risk attaches to where they are.

The internet has none of this. It has sites, and the sites bear no relationship to one another and carry no architectural grammar telling a child where she is or what kind of place this is. Educational content sits beside entertainment beside commerce, and a child on a safe educational platform meets the same interface logic as a child on a safe gaming platform — endless scroll, friend suggestion, notification design — because both are built on the same engagement substrate. The substrate ignores the differentiation that physical geography supplied for nothing.

This is a stronger claim than a category label in a metadata record. A description language says a site is educational. A district commits a registrant to inhabit an environment type and honour its internal grammar — its interaction patterns, its adult-presence patterns, its social-mixing patterns — the way a school building does not permit a slot machine in the cafeteria. Whether that commitment can be specified and enforced is an open question, and among the most interesting the trial could examine.

6.2 Enhancement, not only prohibition

One instrument in the zoning repertoire is routinely overlooked and transposes unusually well. A school zone does not prohibit anything that was not already prohibited. It multiplies the penalty for it. Drug-free school zones, doubled speeding fines, enhanced sentencing within a fixed radius — the offence is unchanged, and proximity to children is an aggravating factor.

That is a different legal move from exclusion, and a milder one. It says not you may not do this, but doing this here costs more. Applied to a namespace, it is the structure of eviction: a violation inside the district costs a registrant standing, admission and the commercial benefit of being there, where the identical conduct outside it would cost nothing.

It is also easy to legislate and hard to characterise as novel censorship, because every jurisdiction already does it and nobody argues that a doubled fine near a school is a restriction on driving.

6.3 What does not transpose

Two things do not carry across, and the argument is stronger for admitting them.

Scarcity, which matters less than it first appears. Land is finite, and a first reading of zoning locates its justification in that finitude. On the health framing set out in Section 3, scarcity was never the ground — externality was. Land is simply where the externalities happened to fall. The question a zoning ordinance actually asks is not whether the

resource is scarce but whether a use imposes a cost on a neighbouring population that cannot avoid bearing it. That question transposes without difficulty: a child cannot avoid a use that is one reach away, and cannot consent to bearing what it costs. The concession is therefore narrower than it looks, though a critic is entitled to press it.

Fixity. A factory does not relocate overnight. A website changes character in an afternoon. Admission at entry is therefore necessary but not sufficient; the regime needs periodic re-evaluation and a route to eviction, which is why the tenancy analogy carries more weight than the permit analogy.

7. What Every Physical Establishment Already Carries

Zoning is one of several overlapping layers within which physical establishments open to the public routinely operate. Which layers apply varies by trade and jurisdiction, but the pattern is familiar, and it is worth setting out because the internet's position is usually described as freedom when it is more accurately described as an exemption nobody granted deliberately.

Consider an establishment offering adult entertainment. It typically operates within four, and none of them is censorship.

- **Zoning.** Whether the premises may exist at that location — distance from schools, places of worship, residential districts.
- **Licensing.** Who may enter, enforced by a person at the door checking identification.
- **Display.** What may face the street, which is a different rule from what may occur inside. The performers are behind a wall, and the wall is a legal requirement rather than an aesthetic choice.
- **Premises liability.** What the operator owes those who come in, because inviting the public onto property creates duties toward them.

The performer's expression is untouched by any of the four. What is regulated is placement, admission, exposure to passers-by, and the operator's duty to those who enter.

7.1 Attractive nuisance

The fourth layer reaches furthest, and its child-specific form is almost too apt. The attractive nuisance doctrine holds a landowner liable for harm to a trespassing child from a hazard likely to draw one — the classic case being an unfenced swimming pool. The child is trespassing and recovers anyway, because children cannot be expected to appreciate the risk.

The remedy has never been to remove the pool. It is to fence it. The doctrine is offered here as illustration rather than authority — it does not govern internet architecture — but it captures a familiar legal intuition: where a lawful thing predictably draws children into harm, the obligation to control access falls on the occupier rather than on the thing to cease existing.

7.2 What a domain has in common with premises

A domain is not real property and a visitor is not, without more, an invitee in the doctrinal sense. The transposition does not require either. A domain performs one function that matters here: it provides an identifiable locus, controlled by an identifiable operator, to which members of the public are invited. A visitor goes to the site rather than the site coming to the visitor — the same directional fact Section 1 describes, seen from the operator's side.

An operator who registers an address, advertises it, optimises for acquisition and sells access is doing what a proprietor does. The vocabulary already concedes it: hosting, visitors, traffic, the master of a domain. Every element of the transaction is commercial, and every equivalent commercial actor in physical space accepts age checking at the door as unremarkable.

The claim that this is all simply free and public consumption was never accurate. Something is being sold, something is being advertised, and members of the public are being invited to an address held by someone who could control who comes in.

7.3 What this does and does not change for age verification

It changes nothing, and the point must be stated plainly because the opposite reading would be damaging.

A child-safe district relieves no operator of any existing obligation. An adult site checks identification for the same reason a licensed premises does: it is selling something age-restricted in a direct transaction, and that duty exists whether or not a children's district exists elsewhere. Any proposal that offered the zone in exchange for relaxed verification would be a deregulation argument wearing protective language, and would deserve the scepticism it received.

What the architecture provides is separation, which age verification currently has to accomplish on its own and cannot. The obligation at the door remains. What changes is that the door is no longer the only line of defence, because children are structurally elsewhere rather than presumed to be everywhere.

8. A Place to Practice

The account so far is negative in form. It defines admission, inspection, and eviction — the wall and its administration. It does not say what the district is for, and the strongest objection to any protected environment turns on exactly that omission.

The objection is that overprotection produces fragile participants. A child who spends every developmental year inside a safe space has learned only safe-space resilience, not the real thing. This is a serious argument and it is made by serious people.

The answer is that the district is not designed to be permanent or total. The closest civic analogy is not a fence but a licensing progression: a learner's permit, then provisional driving, then a full licence — each stage widening exposure as capability grows, none of them permanent.

A child learns to ride a bicycle on a playground before entering traffic, and nobody calls the playground censorship. It is a developmental environment, and its whole purpose is to be left behind.

But the playground also carries the warning. A playground is already designated. It is already zoned for children, marked as theirs, and set apart from the road. And playgrounds still fill with broken glass, discarded bottles, needles and worse — not because the designation was wrong, but because a designation is only a sign on a fence unless something governs what enters and someone maintains it afterwards.

That is the difference between a label and a district, and it is the whole distinction this paper turns on. A children's space needs three things rather than one: a designation, a rule about what may come in, and a continuing obligation to keep it. Admission policy, tenancy, and eviction are those three. Prior attempts at a children's namespace supplied the first and neither of the others.

It also locates the parent correctly. A parent choosing a playground is not being relieved of responsibility; they are exercising it, and the designation is what makes the choice meaningful. The architecture enlists parents rather than replacing them, and a system that gives a parent a real alternative is doing something different from one that hands them another dashboard.

But the parent is not the only party with standing, and a paper that describes the architecture entirely through parents, registrants and regulators has left out the people it is for. Three commitments follow, and they are design requirements rather than sentiments. Children participate in the development of the criteria, not merely as research subjects. A competent adolescent has a route by which their view about their own settings is heard. And the trial measures effects on access, autonomy, participation and privacy alongside reduction in specified harms, because an architecture that made children safer by making them smaller would have failed on its own terms.

The last of those is the one most likely to be quietly dropped, and it should not be. The question is not how safe children can be made. It is whether their rights, protection among them, are better realised inside the district than outside it — and that is a question with a measurable answer that could come back against us.

WHAT THE DISTRICT IS FOR

A bounded space in which a child exercises genuine autonomy — makes real choices, follows real interests, navigates a real internet — but where specified high-severity hazards have been excluded by admission policy rather than left to post-exposure detection. It does not make a child safe from other children, from their own mistakes, or from every category of harm. It removes a defined class of externally supplied hazard so that ranging freely is not reckless.

Failure is instructive rather than ruinous. Resilience is the goal; the district is the training environment, not the destination.

This is also where Article 5 stops being an abstraction. An admission-based district renders the evolving-capabilities principle as a place rather than a dashboard: controls that loosen as the child grows, expressed as an environment the child grows within and eventually out of, rather than a set of prohibitions a parent must perpetually administer.

It follows that the district should be calibrated to the developmental window — most valuable for younger children, deliberately loosening as the child matures, and built to be outgrown. A design that made the space maximally attractive to a seventeen-year-old would have failed at its own purpose.

9. Why This Is Not Censorship

The most consequential objection is that an admission regime is censorship under another name, and a precedent that authoritarian regimes could appropriate. The objection is serious and deserves a structural answer rather than an assurance.

The zoning frame supplies the first part of the answer. Enforcing a land-use code is the administration of an objective rule, not the exercise of editorial judgment. A non-conforming use is removed because it breaches the published code, not because anyone has judged its expressive content unwelcome. The liquor store is not suppressed; it is located elsewhere.

9.1 Six limiting conditions

The architecture is structurally distinct from censorship along six dimensions identified so far. Each names a distinct power a coercive analogue would hold. Whether these are jointly sufficient is treated as an open question in Section 15 rather than asserted here.

Axis	Condition
Compulsion	Participation is voluntary and elected by the household. No person is required to operate within the namespace and none is prevented from operating outside it.
Authorship	The election is private and parental, not governmental. The architecture asks nothing of the state beyond recognition, and functions identically where recognition is withheld.
Externality	No third party is compelled to suppress, remove or modify any expression outside the namespace. The protective property arises from what is admitted, not from what others must remove.
Substitution	The open internet continues unchanged alongside the namespace. This is a parallel district, not a replacement substrate.
Governance	Policy authority is held by an institutionally independent body operating a published process with appeal rights — not by the registry operator, the state, or any single constituency.

Exit	A household may discontinue the election without penalty and without loss of ordinary access to the open internet. A regime that is voluntary at entry can still be coercive if leaving is difficult, punitive, technically obstructed or economically prohibitive.
------	---

These are testable institutional facts rather than preferences. An external observer can verify whether a given deployment satisfies them. A regime satisfying five and failing the sixth fails on a nameable dimension, and that dimension is the material fact a critic would identify.

The architecture is therefore exportable but not appropriable. The primitive — a namespace admitting registrations against a published code — can be replicated anywhere. Replicating it without these conditions produces a different system, one an external observer can identify as such. The defence against authoritarian appropriation is not the hope that bad actors will not copy the primitive; it is that copying the primitive without the conditions does not produce this architecture.

9.2 The legislative option, and what it costs

The six conditions describe an architecture that requires no state action at all. A separate question is what a state might lawfully do once such a district exists.

A legislature may require that a connection supplied to a child be to a governed space. That is not a restriction on publication. It is a restriction on a commercial transaction with a member of a protected class — the structure of every age-restricted goods law there is. The distiller is not censored because the sale to a fourteen-year-old is prohibited; the duty falls on the seller. Applied here, the duty falls on the supplier of the connection, and no publisher is asked to remove, alter or refrain from anything.

Two conditions make that version defensible rather than dangerous. The duty must fall on suppliers rather than on publishers. And the state must not write the admission code. A government that mandates connection to a zone and also defines what is in the zone has performed censorship with additional steps; a government that requires children be connected to a governed space, whose standard is held by an independent international body it does not control, has done something categorically different. This is why the independence of the policy body is structural rather than decorative.

The price should be stated. A state that mandates connection acquires an interest in what the zone contains, and a legislature that requires children to be in the reservoir will eventually have views about the water. The mandate is what makes the architecture reach every child and what makes the standard worth capturing. Those are the same fact, and a jurisdiction adopting the legislative route should adopt the capture safeguards with it.

9.3 Delegation of public power

A distinct objection, and the most serious one a public lawyer would raise. If admission to a space children use is determined by criteria set by a non-governmental international body, then significant authority over access has been located outside the ordinary machinery of democratic accountability and judicial review. The concern is not that the body would act badly. It is that the usual correctives may not reach it.

Three things reduce the force of the objection without disposing of it.

- Participation is voluntary. No person is required to enter the space and none is prevented from operating outside it, so the body's decisions do not determine what anyone may say or read — only what may be admitted to a district a household has chosen.
- Substantive thresholds are set locally. The body maintains the schema; the applicable floor in any jurisdiction is set within that jurisdiction. What is delegated is the format, not the judgment.
- The body's own process is published, its findings documented, and its decisions appealable to a body separate from the one that made them. That is not judicial review, and it is not offered as a substitute for it.

What remains is real and should be conceded. If any jurisdiction moved from voluntary election to mandated connection, the accountability question would sharpen considerably, because the body's decisions would then bear on people who had not chosen them. That is a further reason — beyond the capture risk in Section 9.2 — why this paper defends the voluntary architecture rather than the mandatory one.

9.4 Two directions of enforcement

Describing a registration as a tenancy is doing analytical work, and it invites an objection worth meeting. Domain names are treated in several jurisdictions as intangible property or as contractual rights with property-like features, and suspension or revocation has been litigated on that footing. A regime that removes registrants needs to be clear about the source of its authority.

The answer is contractual rather than statutory, and it runs in two directions rather than one.

Vertically, the registration is a tenancy. Admission is granted on terms accepted by the registrant, those terms incorporate the published code, and removal for breach is enforcement of an agreement rather than exercise of a public power. That is ordinary registry practice: registration agreements already provide for suspension on breach, and disputes are already handled under published policies. Its virtue is speed — a registry can suspend within hours where material requires it, and no court need be troubled. Its weakness is that it is centralised, and therefore capturable: a registry that goes slack, or is leaned on, or simply declines to act, leaves the code unenforced.

Horizontally, the namespace is a scheme of development. English law has a settled structure for exactly this situation. Where land is laid out on a defined plan with reciprocal restrictions, every plot owner may enforce the covenants against every other, regardless of when they bought or in what order. The conditions were set out in *Elliston v Reacher* and refined since: a common vendor, a defined area laid out in lots, restrictions consistent only with a general scheme and intended for the mutual benefit of all lots, and purchasers taking on that footing. The covenant at issue in that case prohibited a hotel, tavern, public house or manufacturing plant on any lot, and nobody has ever described it as censorship.

Each element has a counterpart here. The defined area is the namespace boundary. The general scheme of restrictions is the published code, substantially uniform across every registrant. Notice is the registration agreement incorporating that code as a condition of entry. What follows is that a registrant does not merely owe obligations upward to a registry; each is simultaneously burdened and benefited, and each has standing against the others.

That is the answer to a registry that will not act. A registrant who has invested in a district on the strength of a code that others are ignoring is not dependent on the goodwill of the operator. Distributed standing is slower than suspension and it is not capturable, which is precisely why both are needed: the vertical route handles urgency, the horizontal route handles neglect.

9.5 What this structure is not

It has been argued that treating cyberspace as a place drives propertisation, and that propertisation produces an anticommons — many parties each holding an effective right to exclude, with the result that a resource is used less than it should be. The objection is a serious one and it is directed at exactly the move this paper makes.

A scheme of development is the structural opposite of an anticommons. An anticommons arises from fragmented rights of exclusion over a shared resource. A scheme creates no new exclusion right over anything that already exists: nobody's access to the open internet is diminished, and no participant acquires a veto over any resource outside the district. What it creates is a set of reciprocal obligations within a bounded area, entered voluntarily, where the restriction is the thing each participant joined in order to obtain. Everyone is bound and everyone benefits.

Land-use law is in fact the historical answer to both tragedies. Overuse and externality are managed by regulating use; fragmentation is avoided by not dividing title. That is what is proposed here — conditions of admission to a district, not new property rights over anything.

One design constraint follows, and it is not optional. In *Birdlip v Hunter* the Court of Appeal declined to infer a scheme partly because the vendor had reserved the right to vary the covenants or decline to enforce them. A policy body holding unfettered discretion to waive the code, or to enforce selectively, would undermine the mutuality on which horizontal enforcement depends. The amendment power must therefore be bounded and the enforcement obligation must be real, or the architecture loses the very feature that makes it something other than a landlord's discretion.

And two consequences for drafting. The terms must be published, specific and accepted in advance rather than imposed afterwards, because a contractual remedy depends on what was agreed. And the appeal route must be real, because a process that can remove a registrant's livelihood without independent review will be tested — correctly — in whatever forum is available.

9.6 Precedent containment

A distinct objection operates at the level of precedent: that conditioning a namespace on a content code legitimises the DNS as a general instrument of content control, opening the door to namespaces organised around political or religious suppression under cover of registry policy.

The answer is that the code is tethered, by its own constitutional documents, to a domain that does not generalise. It is confined to objective, categorical, non-viewpoint-based protection of a recognised vulnerable legal class, in a domain where heightened protection is recognised across virtually every legal culture. A code so tethered cannot be extended to justify a namespace organised around ideology, because ideology shares none of the features that justify this one: it is not objective, not categorical, not directed at a recognised vulnerable class, and not the subject of cross-cultural consensus.

The precedent is narrow by construction. An entity invoking this architecture as authority for viewpoint-based restriction would have to abandon the features that make it defensible, at which point it is not invoking this precedent but constructing a different and undefended one.

9.7 Legislative non-cooptation

An objection running the other way is equally serious. Once a child-safe district exists, a legislature might prohibit platforms on the open web from admitting minors at all, converting a voluntary parallel space into the instrument of a mandatory segregation regime.

The answer is a stated commitment rather than a technical safeguard. The legitimacy of the design rests on the continued coexistence of an unchanged open internet; a deployment in which the open web is cleared of minors by statute is not the architecture this paper defends. The foundation commits to oppose, in regulatory comment and as a matter of governance principle, any law using the namespace's existence as grounds for mandatory open-web age segregation. Its lever is withdrawal of endorsement and certification — not the severing of service to a jurisdiction, which would itself become a control point and a capture target.

10. The Code Is a Floor

The code consists of objective, categorical exclusions of the kind a children's protection code characteristically involves, stated as hard rules rather than discretionary judgments — in the same way a residential district categorically excludes heavy industry without adjudicating the merits of any particular factory.

What the architecture fixes is not the list but the character of the code: objective, categorical, tethered to the protection of minors, and revisable through a published process. The substantive content is set and revised by the policy body, and is expected to evolve as consensus and evidence develop, exactly as land-use codes are amended over time.

10.1 Pluralism lives above the floor

Above the floor, the architecture supports a marketplace of curators — in land-use terms, the latitude for each family to develop its own parcel within the code. The code guarantees a common minimum every household can rely on; families then layer higher standards according to their own values, served by competing curators operating on the same machine-readable metadata.

This imposes no fragmented compliance burden on registrants. A registrant satisfies the single predictable code once, at admission, and is tagged accordingly. Curators do not impose new requirements; they provide alternate filtered views of an already-conforming universe. Pluralism therefore lives in the curator layer, where families with differing values are served by differing curators, while the code remains a minimal objective common denominator rather than a contested editorial position.

This answers the concern that a centralised standard would flatten pluralism or marginalise minority communities. The schema contains no field describing the faith, tradition, political alignment or value system of a domain, and will not. Curators read the schema; only the policy body writes it.

10.2 Districts, not one district

The schema is invariant across jurisdictions: the same fields, the same encoding, so that any conforming system reads any national record without loss. The values inside those fields are set locally. This is a technical- interoperability claim, not a content-policy one — the equivalent of ISO 8601, which fixes how a date is written without opinion about which dates matter.

The foundation owns the format. Each jurisdiction owns what goes in it.

11. The Household Election

The unit of consent is the household account, not the individual child. No age verification of individual users is performed, and none is required. A parent elects the protective regime; the household then operates within it. The architecture is closer in form to a family broadband plan than to a moderation regime.

Household election is the entry mechanism, not a claim that every child in a household should receive the same environment. The practical expression must change with evolving capacity: differentiated age bands, attached where possible to distinguishable connections rather than to the household as a whole, and meaningful weight given to the views of older children about their own settings. A design that stopped at the household flag would have implemented convenience rather than Article 5.

11.1 The rights-native anchor

Article 5 of the UN Convention on the Rights of the Child provides the anchor that travels furthest. It locates direction and guidance with the parent as the recognised agent for a child's evolving capacities — the child holds the rights; the parent exercises them on the child's behalf, scaled to the child's developing maturity.

This is the correct framing for three reasons. It is affirmatively codified rather than inferred. It is near-universally ratified. And its evolving-capacities clause supplies an internal limit that a pure parental-authority doctrine lacks: parental direction is calibrated to, and eventually yields to, the child's developing autonomy. An account-level regime implements that clause as a settings surface — controls that evolve with the child rather than controls that are simply imposed.

One consequence of that clause should be built rather than asserted. A single household-wide setting treats a sixteen-year-old and a six-year-old identically, which is not what evolving capacities means and is not defensible under it. The election must therefore be expressible in age bands, with the older bands materially wider, and with a route by which an adolescent's own view is heard as their capacity develops. An architecture that cannot express that difference has implemented parental authority rather than Article 5.

Where the bands fall is not for this paper to decide. Boundaries drawn from convenience rather than from developmental evidence would reproduce, at the level of design, exactly the arbitrariness the architecture criticises elsewhere. The bands are set by the policy body on advice from developmental clinicians and researchers, revised as that evidence develops, and published like any other provision of the code.

Articles 12, 13 and 17 cut in the other direction and must be engaged rather than avoided. The Convention protects the child's right to be heard, to freedom of expression, and to access information from diverse sources. An admission-gated namespace restricts reach by design. General Comment 25 on children's rights in the digital environment is directly on point and is simultaneously this architecture's best authority and its hardest interlocutor. The honest position is that a bounded district which a household may enter and leave, alongside an unchanged open internet, is a narrower interference than a filtering mandate applied to the whole — but the interference is real and the argument must be made rather than assumed.

11.2 The United States is a different jurisdiction, not the model

In the United States the same household election rests on a distinct and well-developed body of doctrine. It recognises substantial parental authority over upbringing and education while also recognising the state's independent protective interest in children; Meyer, Pierce, Prince, Yoder and Troxel mark different points in that relationship rather than a single ascending line, and Prince in particular is a limit rather than an extension.

The interest of parents in the care, custody, and control of their children is perhaps the oldest of the fundamental liberty interests recognized by this Court.

Troxel v. Granville, 530 U.S. 57, 65 (2000).

That footing matters in its own jurisdiction and should not be exported. The United States is the only UN member state that has not ratified the Convention; conversely, Troxel's presumption in favour of the fit parent has no internal analogue to the evolving-capacities limit. The two frameworks reach a similar practical place by different routes and diverge precisely where a competent adolescent dissents from a parental election. That divergence is a genuine research question, not a drafting problem, and it is treated as such.

The United States is also a later phase, and deliberately so. The mechanism this paper describes depends on a national registry able to set second-level policy under its own authority. The American country-code arrangement does not provide that route in the way several other national registries do — an institutional and governance difference rather than a technical impossibility. An American instance therefore requires a generic top-level namespace, which requires an ICANN application, which requires both an available round and a defensible First Amendment posture for admission decisions made at the registry layer. Each of those is tractable and none of them is quick. The sequence follows from the instruments available in each jurisdiction rather than from any judgment about where the need is greatest.

The structural claim is what travels. Whoever the local legal order designates as the child's decision-maker cannot be displaced by the namespace. In the United States that is the parent under Troxel; in Convention states it is the parent as recognised agent under Article 5; elsewhere it may be parent and statutory guardian together. The namespace takes no position on which, and guarantees only that it will not become the decision-maker itself.

12. What Zoning Offers Regulators

The preceding sections argue the architecture from the child's side. This one states what a regulator gets, because the design asks something of regulators and the exchange should be explicit.

The problem regulators currently have. Because the open internet mixes children and adults indiscriminately, every provider must be treated as though children may be present. That single presumption generates much of the compliance weight in the field: duties attach to services never built for children and with no particular reason to attract them, and the same provider carries a different version of those duties in each jurisdiction it can be reached from. The presumption is not a drafting choice. It follows from the absence of structural separation — there is nowhere children are, so they must be assumed to be everywhere.

What a designation changes. Zoning a district for a purpose changes what may be presumed about the parcels outside it. As households with children come to operate inside a designated space, the open web ceases to be where children are structurally presumed to be, and the presumption generating the burden weakens in proportion. That proportion matters, and is stated plainly below.

12.1 Two distinct offers, running in opposite directions

These are separable and would be granted by different instruments. A regulator may take either without the other.

Inside the district: conformance as a defined route to compliance. A registrant admitted under the published code has satisfied an objective, auditable standard at entry and remains subject to inspection and eviction. A regulator may treat that conformance as discharging the child-protection duties the code covers, as a building constructed to an adopted code is not separately litigated for structural adequacy. This is a safe harbour in the conventional sense — a defined route to compliance, not an exemption from it — and what the regulator gains is a population of services whose child-safety posture is machine-readable and verifiable rather than asserted in a policy document.

Outside the district: relief from the presumption, not from any duty. A provider who never enters is among the beneficiaries, relieved of obligations that exist only because no structural separation currently does. The relief is narrow and its limits are the point. It reaches the presumption that children are present. It does not reach any duty arising from what a service sells or does. An operator offering age-restricted goods checks identification at the door whether or not a children's district exists, for the same reason a licensed premises does, and Section 7.3 states that without qualification. What eases is the diffuse general-population duty; every transaction-specific one remains.

12.2 What recognition would require

Recognition does not require new primary legislation in most jurisdictions. The available instruments are ordinary ones, and they differ in what evidence each presupposes.

- **Codes of practice and regulatory guidance.** A regulator with code-making powers may name conformance as one route to compliance without displacing others.
- **Enforcement priority.** A regulator may state that conforming services sit outside its priority enforcement population. This requires no legal change at all and is reversible.
- **Statutory safe harbour.** Where a legislature wishes to go further, a defined affirmative defence for conforming registrants — subject to the conditions in Section 9.2, that the duty fall on suppliers rather than publishers and that the state not author the code.

Each is proportionate to the evidence available at the time. The third should follow a trial rather than precede one.

12.3 The condition, stated rather than buried

The relief is not automatic and does not arrive on the day a district is designated. The presumption of child presence weakens only as households actually move, and until enough do, an open-web provider's exposure is unchanged. That is the same adoption problem Section 13 identifies in every prior attempt, and it applies to this offer as much as to the namespace itself.

Two things follow. A regulator should not be asked to trade present obligations against a future population; the sequence runs the other way, with recognition extended as adoption is demonstrated and measured. And the offer's value is therefore contingent, which is why it is put as a bargain rather than a claim. The trial exists in part to produce the number a regulator would need before striking it.

12.4 What government sheds, and what it keeps

The prevailing approach asks each state to legislate and enforce its own content rules against a borderless internet. That contest is structurally unwinnable in the way filtration is: a national statute binds the publishers a

national court can reach, and much of what a child encounters is published beyond them. It has produced the litigation and proportionality difficulties that statutory approaches have met across jurisdictions.

An admission standard reaches differently. It binds through the registration agreement rather than through territorial jurisdiction, and a registration agreement travels with the registrant wherever the registrant sits. That is a reach no statute has and no regulator can legislate into existence. It is also, for the same reason, a reach that depends entirely on the registrant having chosen to enter — which is the trade.

What government sheds is authorship of the code, not enforcement of it. The relationship is the one Section 5 describes: jurisdictions adopt model building codes; they do not draft fire-resistance ratings. A legislature writing its own would be understood to have exceeded its competence, and the same holds here. Government retains every enforcement power it has today, and gains an objective standard it did not have to write, cannot be accused of authoring, and does not have to defend in court as its own content judgment.

That last point is the durable one. A state defending its own content code defends a viewpoint. A state enforcing an independently authored, published, categorical admission standard is doing what it does when it enforces an adopted building code, and the constitutional posture is entirely different.

13. Why Earlier Attempts Failed

The content-designated namespace is not a new idea. Three attempts are instructive, and they failed in different ways for different reasons — which is more useful than a single diagnosis, because it separates the variables.

Three things must be true for such a namespace to work. There must be authority: some source of legitimacy for the code that the community it governs will accept. There must be distribution: some mechanism by which a household elects and reaches the space. And there must be demand: some reason a registrant would prefer to be inside it. No attempt so far has had more than one.

13.1 kids.us — authority without community or distribution

Created by statute in 2002 as a second-level space under the American country code, with a published child-safety baseline and an operator under contract. The authority was unambiguous: an act of Congress.

It reached roughly 651 registrations and six active sites by 2011 and was suspended in 2012. Nothing changed where a household's traffic resolved, and nothing gave a registrant a reason to move from an address that already had an audience. Statutory authority produced a lawful space nobody entered.

13.2 .xxx — the closest structural precedent, and the clearest warning

This is the case most worth studying, because its architecture is close to the one proposed here. A sponsored top-level domain delegated in 2011, with an independent policy body — the International Foundation for Online Responsibility — responsible for the code, pre-screening of applicants against a charter, a post-registration challenge process, and baseline policies including a prohibition on marketing to minors. Registry operation and policy authority formally separated, as this paper proposes.

It failed on authority, and the failure is precisely located. The policy body was created by the registry it governed, funded by that registry on a per-registration basis, and chaired by the registry's own president until he stepped down following sustained criticism over the conflict. The industry's principal trade association opposed the namespace throughout on the ground that the applicant did not represent the community it claimed to speak for, and ICANN's board declined the application in 2007 substantially on that basis. Delegation followed four years later after an independent review panel found procedural fault in the rejection — a finding about process, not a finding that the community had changed its mind.

The consequences were predictable. Registrations were largely defensive: operators bought names to prevent others holding them and continued publishing elsewhere. Adoption by the community the namespace was built for did not follow, and the space is used mainly for secondary addresses.

The lesson for any independent policy body is exact. Structural separation on paper is not independence. A body created by the operator it governs, funded per transaction by that operator, and chaired by that operator's officers will be described as a subsidiary by the community it claims to serve, and the description will be fair. Independence must be constituted before the commercial relationship exists, not asserted alongside it.

One further detail is worth recording. The same trade association that opposed .xxx observed that a namespace for child-friendly sites would be a more appropriate instrument than a namespace for adult content. The objection was not to designation as such. It was to being designated by someone else, for someone else's purposes.

13.3 .kids — community without distribution or demand

Delegated in 2022 to a not-for-profit foundation, grounded in the UN Convention on the Rights of the Child, with a child-friendly content guideline and an anti-abuse policy binding on registrars. Its community endorsement was genuine, included established child-protection organisations, and satisfied ICANN's community priority evaluation — a demanding test that .xxx did not pass.

And it has not achieved meaningful adoption. This is the most important of the three cases because it isolates the variable. Authority was legitimate and community support was real, and neither was sufficient. Nothing changed where a household's traffic resolved, and nothing gave a registrant a reason to be there rather than somewhere with an audience already.

13.4 What this establishes, and what it does not

Designation, content rules, and even genuine community endorsement do not produce adoption. That is a negative result of some strength, available to anyone proposing a children's namespace, and it should be taken seriously by this paper before it is used against it.

What it does not establish is that the problem is insoluble. None of the three had a distribution mechanism, and none had a registrant incentive beyond the designation itself. Whether supplying those changes the outcome is untested, which is an argument for testing rather than for assuming either result.

13.5 Why authority differs under a country code

Authority is where the country-code route differs most sharply, and it explains the sequence this paper proposes.

A national registry's authority is not derivative. It is not created by the applicant, not funded per registration by the party it constrains, and not vulnerable to the charge that it represents a community which has not consented. It exists independently, is recognised within its jurisdiction, and carries national accountability of its own. A district operated at the second level under that authority begins with a legitimacy .xxx never acquired and kids.us could not transfer.

It also produces the thing a community application needs and cannot manufacture. Where a global namespace is eventually sought, the route carrying real weight is the community one, and its test is whether a delineated community demonstrably exists. .kids passed that test; .xxx did not. A coalition that has operated national trials together, under one schema, with a shared governance body and a record of decisions taken, is a different object from a coalition assembled to support an application — and the difference is visible to anyone assessing it.

The sequence therefore follows from the diagnosis rather than from convenience. Country-code trials supply authority that is not borrowed, and build the community a later application would have to demonstrate. Neither can be created at the point of application, and both are what the earlier attempts lacked.

14. Why the Trial Comes First

The prior failures also determine the sequence. A namespace proposed as a global instrument, defended in the abstract, and delegated before anyone has run one is a repetition of kids.us with better argument.

The order is therefore inverted. A trial at the second level under a participating country-code registry can begin without a new delegation, produce an operating record within a single funding cycle, and answer empirically the questions that are currently answered by assertion: whether admission criteria can be specified precisely enough that independent reviewers reach consistent decisions, what error rate results, what appeal volume an operating registry carries, and what parents do when offered the choice.

Multi-country trials follow, because a schema tested in one language and one legal order has not been tested. The claim that the standard is international rather than Anglophone-with-translation is verified by running it against different languages, different legal definitions, and different cultural thresholds — or it is falsified, which is equally useful.

There is a second reason, and it is structural rather than evidential. Where a global namespace is eventually sought, the route that carries real weight is the community one — an application made on behalf of a delineated community that already exists, with its own membership, practice and history. A community asserted on paper is worth very little against that test. A community that has been operating national trials together, under one schema, with a shared governance body and a record of decisions taken, is a different object entirely.

The trials are therefore not only how the evidence is produced. They are how the community is formed. Each national instance adds a registry, a set of safety organisations, a cohort of trained reviewers and a body of operating practice to something that either exists by then or does not. Nothing else available within the naming system carries comparable weight, and nothing else can be manufactured at the point of application.

Only then does a global namespace become a defensible proposition, because by then the question of whether it works will have an answer that does not depend on the author, and the community on whose behalf it would be sought will be one that can be pointed at.

15. Open Questions

This paper is offered as an argument to be contested rather than a position to be defended. Five questions warrant particular attention.

- **Specifiability.** Can admission criteria be defined precisely enough that independent reviewers applying them to the same material reach the same decision? If inter-rater agreement is poor, the central design assumption fails. This is empirical and answerable.
- **The scarcity gap.** Zoning's justification rests substantially on the externalities of finite land. Namespace is not finite in that way. Does the transposition survive without the scarcity premise, and on what alternative footing?
- **Translation.** Does a classification survive movement between languages and legal orders, or does it fragment into national variants that share a schema and nothing else?
- **The dissenting adolescent.** Where Article 5's evolving-capacities limit and Troxel's parental presumption diverge, which governs, and what does an account-level regime owe a competent minor who wishes to leave?
- **Sufficiency of the conditions.** Are the limiting conditions jointly sufficient to distinguish this architecture from a coercive analogue, or is there a further axis they do not name? Exit was added after the first five proved incomplete, which is reason to expect others.
- **The evidence threshold.** What quality and quantity of evidence should justify a categorical exclusion, particularly where developmental effects are contested? A code excluding on precaution alone is not defensible; a code waiting for settled effect sizes will never act. Where the line sits is a governance question the architecture has not answered.
- **Adoption and incentives.** Can the architecture create sufficient value for publishers and households to overcome the adoption failure demonstrated by earlier children's namespaces? This is the question on which the proposition most plainly turns, and on which there is least evidence.

16. Conclusion

The argument can be stated in one sentence. Regulation of children's online safety has been built on a metaphor that does not describe the thing it regulates, and the discipline that already solved the analogous problem in physical space — deciding what may sit next to what, through published categories and attached permissions — has not been applied.

This is not a proposal for new technology. It is a proposal to apply a settled legal technology to a surface that has never had one. Zoning is a century old, administered daily by a profession with its own statute, case law and literature, and understood across legal cultures as something other than censorship. An admission standard for a namespace is a zoning ordinance rendered machine-readable.

Whether it works is an empirical question, and the answer does not yet exist. That is what the trial is for.

Acknowledgement

The characterisation of the internet as a pile, and the argument that districts must carry internal grammars rather than merely exclude, originate with Professor Christine Satchell of Queensland University of Technology, and are used here with thanks.

About

IFC-SiS — the International Foundation for Child-Safe internet Standards — is an independent non-profit developing an admission-and-metadata standard for child-safe namespaces. Version 1.0 of the vocabulary is published as a draft for public comment. The companion paper, White Paper II: The Admission Layer, treats the technical architecture.

S. Harrison Knudson, J.D.

Director · International Foundation for Child-Safe internet Standards
shknudson@ifcsis.org · +1 202 684 6309